

Homework 3: Supervised machine learning

UIC CS 418, Fall 2023

According to the **Academic Integrity Policy** of this course, all work submitted for grading must be done individually, unless otherwise specified. While we encourage you to talk to your peers and learn from them, this interaction must be superficial with regards to all work submitted for grading. This means you cannot work in teams, you cannot work side-by-side, you cannot submit someone else's work (partial or complete) as your own. In particular, note that you are guilty of academic dishonesty if you extend or receive any kind of unauthorized assistance. Absolutely no transfer of program code between students is permitted (paper or electronic), and you may not solicit code from family, friends, or online forums. Other examples of academic dishonesty include emailing your program to another student, copying-pasting code from the internet, working in a group on a homework assignment, and allowing a tutor, TA, or another individual to write an answer for you. Academic dishonesty is unacceptable, and penalties range from failure to expulsion from the university; cases are handled via the official student conduct process described at <https://dos.uic.edu/conductforstudents.shtml>.

This homework is an individual assignment for all graduate students. Undergraduate students are allowed to discuss in pairs. There will be no extra credit given to undergraduate students who choose to work alone. The pairs of students who choose to work together and submit one homework assignment together still need to abide by the Academic Integrity Policy and not share or receive help from others (except each other).

Due Date

This assignment is due at 11:59pm Tuesday, November 21, 2022. Unfortunately, we cannot extend the deadline.

What to Submit

You need to complete all code and answer all questions denoted by **Q#** (each one is under a bike image) in this notebook.

For this homework, you need to submit both **.ipynb file** and **pdf file**. Both these files together will be considered as submission. Moreover, when submitting the pdf please link the questions in gradeoscope to the appropriate page number in the pdf. **5points** will be deducted if questions arent linked correctly. Make sure to run and display the output of cells related to questions and double-check whether your homework file looks good before and after submission.

```
import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
import nltk
import sklearn
```

```
import string
import re # helps you filter urls
from IPython.display import display, Latex, Markdown
```

Classifying tweets [100 points]

In this problem, you will be analyzing Twitter data extracted using [this](#) api. The data contains tweets posted by the following six Twitter accounts: `realDonaldTrump`, `mike_pence`, `GOP`, `HillaryClinton`, `timkaine`, `TheDemocrats`

For every tweet, there are two pieces of information:

- `screen_name`: the Twitter handle of the user tweeting and
- `text`: the content of the tweet.

The tweets have been divided into two parts - train and test available to you in CSV files. For train, both the `screen_name` and `text` attributes were provided but for test, `screen_name` is hidden.

The overarching goal of the problem is to "predict" the political inclination (Republican/Democratic) of the Twitter user from one of his/her tweets. The ground truth (i.e., true class labels) is determined from the `screen_name` of the tweet as follows

- `realDonaldTrump`, `mike_pence`, `GOP` are Republicans
- `HillaryClinton`, `timkaine`, `TheDemocrats` are Democrats

Thus, this is a binary classification problem.

The problem proceeds in three stages:

- **Text processing (25%)**: We will clean up the raw tweet text using the various functions offered by the `nlTK` package.
- **Feature construction (25%)**: In this part, we will construct bag-of-words feature vectors and training labels from the processed text of tweets and the `screen_name` columns respectively.
- **Classification (50%)**: Using the features derived, we will use `sklearn` package to learn a model which classifies the tweets as desired.

You will use two new python packages in this problem: `nlTK` and `sklearn`, both of which should be available with anaconda. However, NLTK comes with many corpora, toy grammars, trained models, etc, which have to be downloaded manually. This assignment requires NLTK's stopwords list, POS tagger, and WordNetLemmatizer. Install them using:

```
nlTK.download('stopwords')
nlTK.download('wordnet')
nlTK.download('averaged_perceptron_tagger')
nlTK.download('punkt')
nlTK.download('omw-1.4')
```

Verify that the following commands work for you, before moving on.

```
lemmatizer=nltk.stem.wordnet.WordNetLemmatizer()
stopwords=nltk.corpus.stopwords.words('english')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data]   /root/nltk_data...
[nltk_data]   Unzipping taggers/averaged_perceptron_tagger.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package omw-1.4 to /root/nltk_data...
```

Let's begin!

A. Text Processing [25 points]

Your first task is to fill in the following function which processes and tokenizes raw text. The generated list of tokens should meet the following specifications:

1. The tokens must all be in lower case.
2. The tokens should appear in the same order as in the raw text.
3. The tokens must be in their lemmatized form. If a word cannot be lemmatized (i.e, you get an exception), simply catch it and ignore it. These words will not appear in the token list.
4. The tokens must not contain any punctuations. Punctuations should be handled as follows: (a) Apostrophe of the form ' s must be ignored. e.g., She ' s becomes she. (b) Other apostrophes should be omitted. e.g, don ' t becomes dont. (c) Words must be broken at the hyphen and other punctuations.
5. The tokens must not contain any part of a url.

Part of your work is to figure out a logical order to carry out the above operations. You may find `string.punctuation` useful, to get hold of all punctuation symbols. Look for [regular expressions](#) capturing urls in the text. Your tokens must be of type `str`. Use `nltk.word_tokenize()` for tokenization once you have handled punctuation in the manner specified above.

You would want to take a look at the `lemmatize()` function [here](#). In order for `lemmatize()` to give you the root form for any word, you have to provide the context in which you want to lemmatize through the `pos` parameter: `lemmatizer.lemmatize(word, pos=SOMEVALUE)`. The context should be the part of speech (POS) for that word. The good news is you do not have to manually write out the lexical categories for each word because `nltk.pos_tag()` will do this for you. Now you just need to use the results from `pos_tag()` for the `pos` parameter. However, you can notice the POS tag returned from `pos_tag()` is in different format than the expected pos by `lemmatizer`.

pos (Syntactic category): n for noun files, v for verb files, a for adjective files, r for adverb files.

You need to map these pos appropriately. `nltk.help.upenn_tagset()` provides description of each tag returned by `pos_tag()`.

Q1 (15 points):

```
# Convert part of speech tag from nltk.pos_tag to word net compatible
format
# Simple mapping based on first letter of return tag to make grading
consistent
# Everything else will be considered noun 'n'
posMapping = {
# "First_Letter by nltk.pos_tag": "POS_for_lemmatizer"
    "N": 'n',
    "V": 'v',
    "J": 'a',
    "R": 'r'
}
# 14% credits
def process(text, lemmatizer=nltk.stem.wordnet.WordNetLemmatizer()):
    """ Normalizes case and handles punctuation
    Inputs:
        text: str: raw text
        lemmatizer: an instance of a class implementing the
        lemmatize() method
                    (the default argument is of type
nltk.stem.wordnet.WordNetLemmatizer)
    Outputs:
        list(str): tokenized text
    """
    # Text tokenization
    symbols = nltk.word_tokenize(text)

    # Correct punctuation and standardize case.
    symbols = [word.lower() for word in symbols if word.isalnum()]

    # Make a lemmatization of the symbols using part-of-speech tags.
    speechTags = nltk.pos_tag(symbols)

    label_symbols = []
    for word, tag in speechTags:
        # Get the initial letter.
        taggedInitial = tag[0].upper() if tag[0].upper() in posMapping
    else 'N'
        # Translate the initial letter into a another format.
        translateInitial = posMapping[taggedInitial]

        # Utilizing the mapped tag, lemmatize the word.
```

```

        lemmatized_word = lemmatizer.lemmatize(word, translateInitial)
        label_symbols.append(lemmatized_word)

    return label_symbols

```

You can test the above function as follows. Try to make your test strings as exhaustive as possible. Some checks are:

```

# 1 point credit
print(process("I'm doing well! How about you?"))
# ['im', 'do', 'well', 'how', 'about', 'you']

print(process("Education is the ability to listen to almost anything
without losing your temper or your self-confidence."))
# ['education', 'be', 'the', 'ability', 'to', 'listen', 'to',
'almost', 'anything', 'without', 'lose', 'your', 'temper', 'or',
'your', 'self', 'confidence']

print(process("been had done languages cities mice"))
# ['be', 'have', 'do', 'language', 'city', 'mice']

print(process("It's hilarious. Check it out http://t.co/dummyurl"))
# ['it', 'hilarious', 'check', 'it', 'out']

print(process("See it Sunday morning at 8:30a on RTV6 and our RTV6
app. http:..."))
# ['see', 'it', 'sunday', 'morning', 'at', '8', '30a', 'on', 'rtv6',
'and', 'our', 'rtv6', 'app', 'http', '...']
# Here '...' is a special unicode character not in string.punctuation
and it is still present in processed text

['i', 'do', 'well', 'how', 'about', 'you']
['education', 'be', 'the', 'ability', 'to', 'listen', 'to', 'almost',
'anything', 'without', 'lose', 'your', 'temper', 'or', 'your']
['be', 'have', 'do', 'language', 'city', 'mice']
['it', 'hilarious', 'check', 'it', 'out', 'http']
['see', 'it', 'sunday', 'morning', 'at', 'on', 'rtv6', 'and', 'our',
'rtv6', 'app', 'http']

```

Q2 (10 points):

You will now use the `process()` function we implemented to convert the pandas dataframe we just loaded from `tweets_train.csv` file. Your function should be able to handle any data frame which contains a column called `text`. The data frame you return should replace every string in `text` with the result of `process()` and retain all other columns as such. Do not change the order of rows/columns. Before writing `process_all()`, load the data into a `DataFrame` and look at its format:

```
from google.colab import drive
drive.mount('/content/drive')
%cd "/content/drive/MyDrive/CS418_HW3"
```

```
tweets = pd.read_csv("tweets_train.csv", na_filter=False)
display(tweets.head())
```

```
Mounted at /content/drive
/content/drive/MyDrive/CS418_HW3
```

	screen_name	text
0	GOP	RT @GOPconvention: #Oregon votes today. That m...
1	TheDemocrats	RT @DWStweets: The choice for 2016 is clear: W...
2	HillaryClinton	Trump's calling for trillion dollar tax cuts f...
3	HillaryClinton	.@TimKaine's guiding principle: the belief tha...
4	timkaine	Glad the Senate could pass a #THUD / MilCon / ...

```
# 9 points credits
```

```
def process_all(df, lemmatizer=nlk.stem.wordnet.WordNetLemmatizer()):
    """ process all text in the dataframe using process() function.
    Inputs
        df: pd.DataFrame: dataframe containing a column 'text' loaded
        from the CSV file
        lemmatizer: an instance of a class implementing the
        lemmatize() method
        (the default argument is of type
        nlk.stem.wordnet.WordNetLemmatizer)
    Outputs
        pd.DataFrame: dataframe in which the values of text column
        have been changed from str to list(str),
        the output from process() function. Other
        columns are unaffected.
    """
    df['text'] = df['text'].apply(lambda x: process(x, lemmatizer))
    return df
```

```
# test your code
```

```
# 1 point credit
```

```
processed_tweets = process_all(tweets)
print(processed_tweets.head())
```

#	screen_name	text
# 0	GOP	[rt, gopconvention, oregon, vote, today, that,...
# 1	TheDemocrats	[rt, dwstweets, the, choice, for, 2016, be, cl...
# 2	HillaryClinton	[trump, call, for, trillion, dollar, tax, cut,...
# 3	HillaryClinton	[timkaine, guide, principle, the, belief, that...
# 4	timkaine	[glad, the, senate, could, pass, a, thud, milc...

	screen_name	text
0	GOP	[rt, gopconvention, oregon, vote, today, that,...
1	TheDemocrats	[rt, dwstweets, the, choice, for, 2016, be, cl...

```

2 HillaryClinton [trump, call, for, trillion, dollar, tax, cut,...
3 HillaryClinton [timkaine, guide, principle, the, belief, that...
4 timkaine [glad, the, senate, could, pass, a, thud, milc...

```

B. Feature Construction [25 points]

The next step is to derive feature vectors from the tokenized tweets. In this section, you will be constructing a bag-of-words TF-IDF feature vector. But before that, as you may have guessed, the number of possible words is prohibitively large and not all of them may be useful for our classification task. We need to determine which words to retain, and which to omit. A common heuristic is to construct a frequency distribution of words in the corpus and prune out the head and tail of the distribution. The intuition of the above operation is as follows. Very common words (i.e. stopwords) add almost no information regarding similarity of two pieces of text. Similarly with very rare words. NLTK has a list of in-built stop words which is a good substitute for head of the distribution. We will consider a word rare if it occurs only in a single document (row) in whole of `tweets_train.csv`.

Q3 (15 points):

Construct a sparse matrix of features for each tweet with the help of `sklearn.feature_extraction.text.TfidfVectorizer` (documentation [here](#)). You need to pass a parameter `min_df=2` to filter out the words occurring only in one document in the whole training set. Remember to ignore the stop words as well. You must leave other optional parameters (e.g., `vocab`, `norm`, etc) at their default values. But you may need to use parameters like `lowercase` and `tokenizer` to handle `processed_tweets` that is a list of tokens (not raw text).

```

# 14 points credits
from sklearn.feature_extraction.text import TfidfVectorizer
def create_features(processed_tweets, stop_words):
    """ creates the feature matrix using the processed tweet text
    Inputs:
        processed_tweets: pd.DataFrame: processed tweets read from
        train/test csv file, containing the column 'text'
        stop_words: list(str): stop_words by nltk stopwords (after
        processing)
    Outputs:
        sklearn.feature_extraction.text.TfidfVectorizer: the
        TfidfVectorizer object used
        we need this to tranform test tweets in the same way as
        train tweets
        scipy.sparse.csr.csr_matrix: sparse bag-of-words TF-IDF
        feature matrix
    """
    processed_tweets['text'] = processed_tweets['text'].apply(lambda
x: ' '.join(x))
    stop_words = list( stop_words)
    tokenizer = lambda x: x.split() # Define a lambda function to

```

```

tokenize the tweet text
    vectorize_tweets = TfidfVectorizer(lowercase=False,
stop_words=stop_words, min_df=2, tokenizer=tokenizer)
    feature_matrix =
vectorize_tweets.fit_transform(processed_tweets['text'])
    return vectorize_tweets, feature_matrix

# execute this code
# 1 point credit
# It is recommended to process stopwords according to our data
cleaning rules
tweets = pd.read_csv("tweets_train.csv", na_filter=False)
processed_tweets = process_all(tweets)
processed_stopwords = set(np.concatenate([process(word) for word in
stopwords]))

(tfidf, X) = create_features(processed_tweets, processed_stopwords)
# Ignore warning
tfidf, X
# Output (should be similar):
# (TfidfVectorizer(lowercase=False, min_df=2,
#                  stop_words={'a', 'about', 'above', 'after',
# 'again', 'against',
#                               'ain', 'all', 'an', 'and', 'any',
# 'aren', 'arent',
#                               'at', 'be', 'because', 'before',
# 'below', 'between',
#                               'both', 'but', 'by', 'can', 'couldn',
# 'couldn't',
#                               'd', 'didn', 'didn't', 'do',
# 'doesn', ...},
#                  tokenizer=<function
create_features.<locals>.<lambda> at 0x7fd4002a6700>),
# <17298x8114 sparse matrix of type '<class 'numpy.float64'>'
#   with 170355 stored elements in Compressed Sparse Row format>)

/usr/local/lib/python3.10/dist-packages/sklearn/feature_extraction/
text.py:528: UserWarning: The parameter 'token_pattern' will not be
used since 'tokenizer' is not None'
    warnings.warn(

(TfidfVectorizer(lowercase=False, min_df=2,
                  stop_words=['until', 're', 'll', 'our', 'their',
'for', 'these',
                              'need', 'won', 'off', 'other', 'hasn',
'herself',
                              'by', 'no', 'of', 'down', 'his', 'needn',
'shan',
                              'nor', 'because', 'with', 'your', 'from',
'if',

```



```

        'but', 'him', 'have', 'i', ...],
        tokenizer=<function create_features.<locals>.<lambda>
at 0x7fc124ec0670>),
<17298x7558 sparse matrix of type '<class 'numpy.float64'>'
with 166256 stored elements in Compressed Sparse Row format>)

```

Q4 (10%):

Also for each tweet, assign a class label (0 or 1) using its `screen_name`. Use 0 forrealDonaldTrump, mike_pence, GOP and 1 for the rest.

```

# 9 point credits
from sklearn.preprocessing import LabelEncoder
def create_labels(processed_tweets):
    """ creates the class labels from screen_name
    Inputs:
        processed_tweets: pd.DataFrame: tweets read from train file,
        containing the column 'screen_name'
    Outputs:
        numpy.ndarray(int): dense binary numpy array of class labels
    """
    labelEncoder = LabelEncoder()
    labelResult =
    labelEncoder.fit_transform(processed_tweets['screen_name'])

    return labelResult

# execute this code
# 1 point credit
y = create_labels(processed_tweets)
y
# 0          0
# 1          1
# 2          1
# 3          1
# 4          1
#           ..
# 17293      0
# 17294      0
# 17295      0
# 17296      1
# 17297      0
# Name: screen_name, Length: 17298, dtype: int32
array([0, 2, 1, ..., 3, 1, 0])

```

C. Classification [50 points]

And finally, we are ready to put things together and learn a model for the classification of tweets. The classifier you will be using is `sklearn.svm.SVC` (Support Vector Machine).

At the heart of SVMs is the concept of kernel functions, which determines how the similarity/distance between two data points is computed. `sklearn`'s SVM provides four kernel functions: `linear`, `poly`, `rbf`, `sigmoid` (details [here](#)) but you can also implement your own distance function and pass it as an argument to the classifier.

Through the various functions you implement in this part, you will be able to learn a classifier, score a classifier based on how well it performs, use it for prediction tasks and compare it to a baseline.

Specifically, you will carry out the following tasks (Q5-9) in order:

1. Implement and evaluate a simple baseline classifier `MajorityLabelClassifier`.
2. Implement the `learn_classifier()` function assuming `kernel` is always one of `{linear, poly, rbf, sigmoid}`.
3. Implement the `evaluate_classifier()` function which scores a classifier based on accuracy of a given dataset.
4. Implement `best_model_selection()` to perform cross-validation by calling `learn_classifier()` and `evaluate_classifier()` for different folds and determine which of the four kernels performs the best.
5. Go back to `learn_classifier()` and fill in the best kernel.

Q5 (10 points):

To determine whether your classifier is performing well, you need to compare it to a baseline classifier. A baseline is generally a simple or trivial classifier and your classifier should beat the baseline in terms of a performance measure such as accuracy. Implement a classifier called `MajorityLabelClassifier` that always predicts the class equal to **mode** of the labels (i.e., the most frequent label) in training data. Part of the code is done for you. Implement the `fit` and `predict` methods. Initialize your classifier appropriately.

```
from sklearn.feature_extraction.text import TfidfVectorizer
from scipy.stats import mode
from sklearn.svm import SVC
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score

# Skeleton of MajorityLabelClassifier is consistent with other sklearn
# classifiers
# 8% credits
class MajorityLabelClassifier():
    """
    A classifier that predicts the mode of training labels
    """
    def __init__(self):
```

```

    """
    Initialize your parameter here
    """
    self.mode_label = None

def fit(self, X, y):
    """
    Implement fit by taking training data X and their labels y and
    finding the mode of y
    i.e. store your learned parameter
    """
    self.mode_label = mode(y).mode.item()

def predict(self, X):
    """
    Implement to give the mode of training labels as a prediction
    for each data instance in X
    return labels
    """
    if X is None:
        return np.array([self.mode_label])
# Initialize your classifier
baselineClf = MajorityLabelClassifier()

# Fit the classifier to the training data
baselineClf.fit(None, y)

# Predict the labels for the training data
predicted_labels = baselineClf.predict(None)

# Calculate the accuracy by comparing the predicted labels with the
true labels
accuracy = np.mean(predicted_labels == y)

# Print the training accuracy
print("Training Accuracy: ", accuracy)
Training Accuracy:  0.16834316105908198

```

Q6 (10 points):

Implement the `learn_classifier()` function assuming `kernel` is always one of `{linear, poly, rbf, sigmoid}`. Stick to default values for any other optional parameters.

Hint: Check <https://scikit-learn.org/stable/modules/svm.html#svm-kernels> on how to use `sklearn.svm.SVC`

```

# 9 points credits
def learn_classifier(X_train, y_train, kernel):
    """ learns a classifier from the input features and labels using

```

```

the kernel function supplied
Inputs:
    X_train: scipy.sparse.csr.csr_matrix: sparse matrix of
features, output of create_features()
    y_train: numpy.ndarray(int): dense binary vector of class
labels, output of create_labels()
    kernel: str: kernel function to be used with classifier.
[linear|poly|rbf|sigmoid]
Outputs:
    sklearn.svm.SVC: classifier learnt from data
"""
classifier = SVC(kernel=kernel)
# Fit the classifier to the training data.
classifier.fit(X_train, y_train)
return classifier

# execute code
# 1 point credit
classifier = learn_classifier(X, y, 'linear')

```

Q7 (10 points):

Now that we know how to learn a classifier, the next step is to evaluate it, ie., characterize how good its classification performance is. This step is necessary to select the best model among a given set of models, or even tune hyperparameters for a given model.

There are two questions that should now come to your mind:

1. What data to use?

- **Validation Data:** The data used to evaluate a classifier is called **validation data** (or hold-out data), and it is usually different from the data used for training. The model or hyperparameter with the best performance in the held out data is chosen. This approach is relatively fast and simple but vulnerable to biases found in validation set.
- **Cross-validation:** This approach divides the dataset in k groups (so, called k -fold cross-validation). One of group is used as test set for evaluation and other groups as training set. The model or hyperparameter with the best average performance across all k folds is chosen. For this question you will perform 4-fold cross validation to determine the best kernel. We will keep all other hyperparameters default for now. This approach provides robustness toward biasness in validation set. However, it takes more time.

2. And what metric?

There are several evaluation measures available in the literature (e.g., accuracy, precision, recall, F-1, etc) and different fields have different preferences for specific metrics due to different goals. We will go with accuracy. According to wiki, **accuracy** of a classifier measures the fraction of all data points that are correctly classified by it; it is the ratio of the number of correct classifications to the total number of (correct and incorrect) classifications. `sklearn.metrics` provides a number of performance metrics.

Now, implement the following function.

```
# 9 points credits
def evaluate_classifier(classifier, X_validation, y_validation):
    """ evaluates a classifier based on a supplied validation data
    Inputs:
        classifier: sklearn.svm.classes.SVC: classifier to evaluate
        X_validation: scipy.sparse.csr.csr_matrix: sparse matrix of
        features
        y_validation: numpy.ndarray(int): dense binary vector of class
        labels
    Outputs:
        double: accuracy of classifier on the validation data
    """
    aoc = classifier.score(X_validation, y_validation)
    return aoc

# test your code by evaluating the accuracy on the training data
# 1 point credit
accuracy = evaluate_classifier(classifier, X, y)
print(accuracy)
# should give around 0.9545034107989363
0.9146722164412071
```

Q8 (10 points):

Now it is time to decide which kernel works best by using the cross-validation technique. Write code to split the training data into 4-folds (75% training and 25% validation) by shuffling randomly. For each kernel, record the average accuracy for all folds and determine the best classifier. Since our dataset is balanced (both classes are in almost equal proportion), `sklearn.model_selection.KFold` can be used for cross-validation.

```
kf = sklearn.model_selection.KFold(n_splits=4, random_state=1,
shuffle=True)
kf

KFold(n_splits=4, random_state=1, shuffle=True)
```

Then use the following code to determine which classifier is the best.

```
# 10 points credits
def best_model_selection(kf, X, y):
    """
    Select the kernel giving best results using k-fold cross-
    validation.
    Other parameters should be left default.
    Input:
        kf (sklearn.model_selection.KFold): kf object defined above
    """
```

```

X (scipy.sparse.csr.csr_matrix): training data
y (array(int)): training labels
Return:
best_kernel (string)
"""
kernels = ['linear', 'rbf', 'poly', 'sigmoid']
avg_accuracies = []
for kernel in ['linear', 'rbf', 'poly', 'sigmoid']:
    # Use the documentation of KFold cross-validation to split ..
    # training data and test data from create_features() and
create_labels()
    # call learn_classifier() using training split of kth fold
    # evaluate on the test split of kth fold
    # record avg accuracies and determine best model (kernel)
#return best kernel as string
    accuracies = []
    for train_index, test_index in kf.split(X):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        clf = SVC(kernel=kernel)
        clf.fit(X_train, y_train)
        accuracy = clf.score(X_test, y_test)
        accuracies.append(accuracy)

    avg_accuracy = np.mean(accuracies)
    avg_accuracies.append(avg_accuracy)

best_kernel_index = np.argmax(avg_accuracies)
best_kernel = kernels[best_kernel_index]

return best_kernel
#Test your code
best_kernel = best_model_selection(kf, X, y)
best_kernel
{"type": "string"}

```

Q9 (10 points)

We're almost done! It's time to write a nice little wrapper function that will use our model to classify unlabeled tweets from tweets_test.csv file.

```

# 8points credits
def classify_tweets(tfidf, classifier, unlabeled_tweets):
    """ predicts class labels for raw tweet text
    Inputs:
        tfidf: sklearn.feature_extraction.text.TfidfVectorizer: the
TfidfVectorizer object used on training data

```

```

        classifier: sklearn.svm.SVC: classifier learned
        unlabeled_tweets: pd.DataFrame: tweets read from
tweets_test.csv
    Outputs:
        numpy.ndarray(int): dense binary vector of class labels for
unlabeled tweets
    """
    # Preprocess the unlabeled tweets
    unlabeled_features = tfidf.transform(unlabeled_tweets['text'])

    # Make predictions using the classifier
    predictions = classifier.predict(unlabeled_features)

    # Return the predicted class labels
    return predictions

# Fill in best classifier in your function and re-train your
classifier using all training data
# Get predictions for unlabelled test data
# 2 points credits
classifier = learn_classifier(X, y, best_kernel)
unlabeled_tweets = pd.read_csv("tweets_test.csv", na_filter=False)
y_pred = classify_tweets(tfidf, classifier, unlabeled_tweets)

```

Did your SVM classifier perform better than the baseline (while evaluating with training data)? Explain in 1-2 sentences how you reached this conclusion.

YOUR ANSWER HERE

Yes, we found the SVM classifier did perform much better than the baseline since it produce the most accurate guesses. Therefore, we assume the SVM model was able to learn patterns in the training data that resulted in more accurate predictions compared to a simple baseline approach. This shows that the SVM classifier is a more effective model for the given task and dataset.